

CHL Lab

## 10/27/17-11/01/17: Checking out TACC platforms

Rosalind Xu 18'

login: ssh -l rosjxu [login.xsede.org](http://login.xsede.org)

pw: Hc

see a list of platforms: xsede-gsissh-hosts

to switch into a resource: gsissh stampede2

OR: ssh directly to

[rosjxu@stampede2.tacc.utexas.edu](mailto:rosjxu@stampede2.tacc.utexas.edu)

(pw: Hc)

### Three Directories:

\$HOME

\$WORK

\$SCRATCH

All computations should be conducted in \$SCRATCH (files will be purged if access time is more than 10 days old). Files can then be moved to \$HOME or \$WORK for permanent storage.

To copy from/to \$WORK:

scp ./myfile rosjxu@stampede2.tacc.utexas.edu:\\$WORK/

Use rsync to update a large data set (instead of copying and pasting every single small file):

Syncing ./mybigfile with rosjxu@stampede2.tacc.utexas.edu:\\$WORK/data/  
mybigfile:

rsync -avtr mybigfile rosjxu@stampede2.tacc.utexas.edu:\\$WORK/data

OR:

rsync -avtr mybigfile/\* rosjxu@stampede2.tacc.utexas.edu:\\$WORK/data/  
mybigfile

DO not move files too frequently between \$HOME / \$WORK and \$SCRATCH,  
especially if that file is large!

## Sanity Check:

module load sanitytool

sanitycheck

## Good Citizenship:

*Log in nodes:* A login node is a good place to *edit and manage files, initiate file transfers, compile code, submit new jobs, and track existing jobs.*

1. Don't run research applications (R, MATLAB, etc.)
2. Don't launch too many simultaneous processes: `make -j 16` (compile on 16 cores) is rude
3. Don't run anything resource intensive

## Shared Lustre file systems:

1. Stripe the receiving directory before transferring large files (a few hundred GB) to that directory
2. Don't run jobs in \$HOME. \$HOME is NOT for parallel jobs!
3. Run I/O intensive jobs in \$SCRATCH rather than \$WORK (remember your core load < 100 % problems?)
  1. If you stress \$WORK, you affect every user on every TACC system!
4. If you suspect your workflow is I/O intensive, don't submit a bunch of simultaneous jobs.
  1. Writing restart/snapshot files can stress the system; avoid doing so too frequently.
5. Watch your file quotas.
  1. If you are near quota in \$WORK and your script keeps trying and failing to write to \$WORK, you will stress the system
  2. If you are near quota in \$HOME, jobs run on \$HOME, \$WORK, or \$SCRATCH may fail, since all jobs write to \$HOME/.slurm
6. Avoid opening and closing files repeatedly in tight loops.
7. Avoid storing many small files in a single directory, and avoid workflows that require many small files.
  1. A few hundred small files is fine, but tens of thousands is too many.
  2. Group small files in separate directories of manageable size.

## Internal and external networks:

1. Avoid too many simultaneous file transfers.
  1. Two or three concurrent scp sessions is fine; twenty is not
2. Avoid recursive file transfers (`scp -r`), especially one involving many small files!
  1. Create a tar archive before transfers: `tar -cvzf [filename].tar.gz [filename]`

### Submitting jobs:

1. Do not ask for more time than you need.
  1. The scheduler will have an easier time to find you a 2 h slot instead of a 48 h slot; which means shorter waiting time for you and everyone else.
2. Test your submission scripts.
  1. Make sure everything works on 2 nodes before you try 200
  2. Debug the submission code with 5-minute jobs: "hello world" codes; one-liners like "ibrun hostname".

### Tar:

tar -cvzf [filename].tar.gz [filename]   —> compress  
tar -xvzf [filename].tar.gz               —> decompress

### Running Jobs:

*Job Accounting:*  $SU = \#nodes * \#wall-clock-hours * job-multiplier$

### Slurm Partitions:

Queue	Node Type	Max Nodes, (assoc'd cores per job*)	Max Duration	Max jobs in queue*	Charge (per node-hour)	Configuration (memory-cluster mode)***
development	KNL	8 nodes (544 cores)*	2 hrs	1*	1 Service Unit (SU)	cache-quadrant
normal	KNL	256 nodes (17,408 cores)*	48 hrs	50*	1 SU	cache-quadrant
large**	KNL	2048 nodes (139,264 cores)*	48 hrs	5*	1 SU	cache-quadrant
flat-quadrant	KNL	32 nodes (2,176 cores)*	48 hrs	4*	1 SU	flat-quadrant
flat-snc4	KNL	12 nodes (816 cores)*	48 hrs	1*	1 SU	flat-SNC4

### Slurm Commands:

sbatch                                       —> job submission  
sbatch --dependency=afterok:[JOBID] [new-job-script]  
  —> start new-job-script after [JOBID]

completes successfully

sinfo -o "%18P %8a %16F"   —> monitoring node availability  
squeue -u rosju             —> monitoring job status  
showq -u                   —> alternative way to monitor job status  
squeue --start -j [JOBID]   —> estimate start time for a job  
scancel [JOBID]             —> cancel job

scontrol show job=[JOBID] —> detailed info on job  
sacct -u rosjxu --starttime 2017-08-01 --  
format=JobID,JobName,MaxRSS,Elapsed  
—> statistics on completed jobs started on or  
after 2017-08-01

#### *Environmental variables:*

Do not use the Slurm "--export" option to manage your job's environment: doing so can interfere with the way the system propagates the inherited environment. Instead, use export env-var=\$val to set environmental variables before submitting a job.

#### *Types of Jobs:*

1. Single serial jobs
  1. Serial jobs require 1 task only per job
  2. -N=1 -n=1
2. Single MPI jobs: ibrun [job-script]
  1. If there is little rush in time, it is recommended to run on single node ( -N=1 )
  2. The max number of mpi processes ( -n ) per node is 64-68. Start small (say 32), then increase gradually
3. Single OpenMP jobs
4. More than one serial application in the same job: module *launcher*
  1. module load launcher
  2. Read \$TACC\_LAUNCHER\_DIR/README
  3. With launcher, multiple jobs can be submitted to the same node
  4. More processes on same node >> using more nodes
  5. Do scaling to find out optimal # of jobs per node
  6. -N=nNodes -n=nNodes\*optimal#Jobs
5. Interactive sessions: idev or srun
  1. idev -p normal -N 2 -n 8 -m 150 # normal queue, 2 nodes, 8 total tasks, 150 minutes
  2. srun --pty -N 2 -n 8 -t 2:30:00 -p normal /bin/bash -l # same conditions as above

## **GROMACS**

gmx —> run gmx features  
ibrun mdrun\_mpi —> run mdrun with MPI parallelization

**SoIEFP:**

python run\_biomol.py —> run biomol

python run\_small.py —> run small molecules

slv\_calc-expres, slv\_calc-ftir —> run slv modules